

## Project Notes

---

### Table Creation

- **UNSIGNED** prevents accidental inserts of negative values (EVERY TABLE)
- **DECIMAL(7,2)** limits price to 99999.99 (too small for some products) (PRODUCT, ORDER\_PRODUCT)
- Mix **0** and **FALSE** (PRODUCT, USER)
- Most university systems run: (PRODUCT, ORDER\_PRODUCT)
  - MySQL 5.7
  - OR MySQL 8.0.x in compatibility mode  
→ CHECK constraints do **nothing**.
- **USER** is a reserved word in MySQL (USER)
- CHAR(13) UNIQUE **with dash**: YYYYMMDD-NNNN → 13 chars
- **Name, Password** are likely to collide with functions (PASSWORD() and NAME() exist in MySQL).
- Better avoid backticks like '**USER**' because:
  - you must ALWAYS escape them
  - queries become harder to read
  - errors are more likely
  - tools might break
- Specification: "Payment reference ... 64 character long string.", so PaymentReference VARCHAR(**64**) NOT NULL

---

### Queries

#### Query 3

- Returns extra fields (**BasePriceNet**, **DiscountPercent**) and gives the price alias a different, less spec-aligned name (**FinalPriceWithVAT** instead of something like **CurrentRetailPrice**).

- DiscountPercent is **NOT NULL** with default 0.

That means **IFNULL(P.DiscountPercent, 0)** will always be the same as just P.DiscountPercent with the current design and data.

#### Query 4

- Forgot to select **Short Description** and **Current Retail Price**.

#### Query 5

- Needs **COALESCE(ROUND(AVG(UP.Stars), 2), 0)** AS AverageRating in order to return 0 average rating when a product has no reviews

#### Query 7

- The 30-day filter is correct, but it makes no visible difference because all our Orders have recent NOW() dates, so none are older than 30 days — the filter becomes meaningful only if some OrderDate values are more than 30 days old.

#### No order dates older than 30 days old

ProductID	Title	TotalUnitsSold
33	Smart RGB Bulb E27	5
10	BassTube Waterproof	4
3	EduBook Air 13	3
11	ErgoMaster 3000	2
7	OLED Cinema 65"	2
30	RCA to Aux Adapter	2
17	Active Stylus Gen 2	1
12	Architect Desk Lamp	1
21	Cast Iron Skillet	1
18	Folio Keyboard Case	1

#### Order dates older than 30 days old

ProductID	Title	TotalUnitsSold
3	EduBook Air 13	3
10	BassTube Waterproof	2
30	RCA to Aux Adapter	2
12	Architect Desk Lamp	1
21	Cast Iron Skillet	1
11	ErgoMaster 3000	1
15	GaN Fast Charger 100W	1
2	Hydra Gaming Tower	1
6	Nordic Commuter Backpack	1
7	OLED Cinema 65"	1

Take for example the **BassTubeWaterproof** product. Before the change on dates the company had sold 4 units **BassTubeWaterproof** of in the last 30 days but after the change the company sold 2 units **BassTubeWaterproof**.

# Queries Analysis

## Query 1:

'-> Index lookup on Department using ParentDepartmentID (ParentDepartmentID = NULL), with index condition: (department.ParentDepartmentID is null) (cost=0.35 rows=1)

- **EXPLAIN** shows that MySQL uses an **index lookup** on the Department table via the ParentDepartmentID index.
- The query efficiently retrieves rows using this existing index, so **no further optimization or additional indexing is needed**.

## Query 2:

'-> Sort: department.Title (cost=0.7 rows=2)

-> Index lookup on Department using ParentDepartmentID (ParentDepartmentID = 1) (cost=0.7 rows=2)

- MySQL performs an **index lookup** using the existing ParentDepartmentID index to retrieve the matching departments efficiently.
- A small **sort operation** on Title occurs due to ORDER BY, but with only two rows this cost is negligible.
- The query is already optimized, and **no additional indexing is required**.

## Query 3:

'-> Sort: product.Title (cost=3.55 rows=33)

-> Filter: (product.IsFeatured = true) (cost=3.55 rows=33)

-> Table scan on Product (cost=3.55 rows=33)

- MySQL scans **all 33 rows** of Product and then filters by IsFeatured.
- This is a **full table scan** (type = ALL), so the query is not optimized.

'-> Index lookup on Product using idx\_product\_isfeatured\_title (IsFeatured = true) (cost=1.75 rows=10)

- After adding the composite index idx\_product\_isfeatured\_title (IsFeatured, Title), **EXPLAIN** shows an **index lookup** instead of a full table scan.

- MySQL now jumps directly to rows where `IsFeatured = TRUE` and reads them already ordered by `Title`.
- Rows examined drop from **33** to **10**, improving performance under the project's "rows examined" metric.

#### **Query 4:**

```
'-> Sort: p.Title
-> Table scan on <temporary> (cost=5.42..7.54 rows=5.89)
-> Temporary table with deduplication (cost=4.98..4.98 rows=5.89)
-> Nested loop inner join (cost=3.62 rows=5.89)
-> Nested loop inner join (cost=1.56 rows=5.89)
-> Covering index lookup on pk_base using PRIMARY (ProductID =
(@ProductID)) (cost=0.45 rows=2)
-> Filter: (pk_other.ProductID <> <cache>(@ProductID)) (cost=0.398
rows=2.94)
-> Covering index lookup on pk_other using KeywordID (KeywordID =
pk_base.KeywordID) (cost=0.398 rows=3.06)
-> Single-row index lookup on p using PRIMARY (ProductID =
pk_other.ProductID) (cost=0.267 rows=1)

'

```

- MySQL performs **covering index lookups** on `Product_Keyword` for both the base product and related products, using existing B-tree indexes on `ProductID` and `KeywordID`.
- For each related product ID, the database performs a **single-row primary key lookup** on `Product`, which is the most efficient access method.
- A **temporary table with deduplication** is created due to the `DISTINCT` clause, followed by a small sort on `Title` required by `ORDER BY`.
- No full table scans occur on the base tables, and the execution plan is already efficient.

## Query 5:

```
'-> Sort: p.Title  
-> Stream results (cost=1.86 rows=1.41)  
-> Group aggregate: avg(up.Stars) (cost=1.86 rows=1.41)  
-> Nested loop left join (cost=1.4 rows=2)  
    -> Index lookup on p using DepartmentID (DepartmentID =  
(@DepartmentID)) (cost=0.7 rows=2)  
        -> Index lookup on up using PRIMARY (ProductID = p.ProductID)  
(cost=0.3 rows=1)
```

- **EXPLAIN** shows an **index lookup** on `Product.DepartmentID`, allowing MySQL to retrieve only products belonging to the selected department instead of scanning the whole table.
- For each product, MySQL performs a **single-row primary key lookup** on `User_Product` to compute the average rating.
- A small sort on `p.Title` is required due to `ORDER BY`, but the number of rows involved is minimal.
- No full table scans occur, and existing indexes already support the query efficiently, so **no additional indexing is needed**.

## Query 6:

```
'-> Sort: product.DiscountPercent DESC, product.Title (cost=3.55 rows=33)  
-> Filter: (product.DiscountPercent <> 0.00) (cost=3.55 rows=33)  
-> Table scan on Product (cost=3.55 rows=33)
```

- **EXPLAIN** shows a **full table scan** on `Product`, examining all 33 rows before filtering products with `DiscountPercent <> 0`.
- MySQL then performs a separate **sort operation** on `(DiscountPercent DESC, Title)` because no index supports this ordering.
- This makes the query a strong candidate for optimization by adding a composite index.

- '-> Sort: product.DiscountPercent DESC, product.Title (cost=2.37 rows=10)
- > Filter: (product.DiscountPercent <> 0.00) (cost=2.37 rows=10)
- > Covering index range scan on Product using idx\_product\_discount\_title over (DiscountPercent < 0.00) OR (0.00 < DiscountPercent) (cost=2.37 rows=10)
- '
- After adding the composite index idx\_product\_discount\_title (DiscountPercent, Title), **EXPLAIN** shows a **covering index range scan** instead of a full table scan.
  - MySQL now examines only the rows that match the discount condition (**DiscountPercent <> 0**), reducing the estimated rows examined from **33 to 10**.
  - Although a small sort step remains for ORDER BY, it now operates on a much smaller filtered set, improving query performance under the project's "rows examined" metric.

## Query 7:

- '-> Limit: 10 row(s)
- > Sort: TotalUnitsSold DESC, p.Title, limit input to 10 row(s) per chunk
- > Table scan on <temporary>
- > Aggregate using temporary table
- > Nested loop inner join (cost=3.98 rows=4.67)
- > Nested loop inner join (cost=2.35 rows=4.67)
- > Filter: (o.OrderDate >= <cache>((curdate() - interval 30 day)))  
(cost=0.717 rows=4.67)
- > Table scan on o (cost=0.717 rows=14)
- > Index lookup on op using PRIMARY (OrderID = o.OrderID)  
(cost=0.271 rows=1)
- > Single-row index lookup on p using PRIMARY (ProductID = op.ProductID) (cost=0.271 rows=1)
- '

- **EXPLAIN** shows a **full table scan** on Orders before applying the date filter, so all rows are examined.
- The joins on Order\_Product and Product use **index lookups**, so those parts are already efficient.
- Grouping, sorting, and LIMIT are expected; the main optimization opportunity is to **index OrderDate** to reduce examined rows.

## Results with the original dataset (no older orders)

```
'-> Limit: 10 row(s)
 -> Sort: TotalUnitsSold DESC, p.Title, limit input to 10 row(s) per chunk
      -> Table scan on <temporary>
          -> Aggregate using temporary table
              -> Nested loop inner join (cost=11.5 rows=14)
                  -> Nested loop inner join (cost=6.55 rows=14)
                      -> Filter: (o.OrderDate >= <cache>((curdate() - interval 30 day)))
                          (cost=1.65 rows=14)
                          -> Covering index scan on o using idx_orders_orderdate_orderid
                              (cost=1.65 rows=14)
                              -> Index lookup on op using PRIMARY (OrderID = o.OrderID)
                                  (cost=0.257 rows=1)
                                  -> Single-row index lookup on p using PRIMARY (ProductID =
                                      op.ProductID) (cost=0.257 rows=1)
```

- All 14 orders were within the last 30 days, so the date filter matched every row.
- **EXPLAIN** before index: **full table scan**, 14 rows examined.
- **EXPLAIN** after index: **covering index scan**, still 14 rows examined.
- No improvement was visible because the filter could not exclude any rows.

## Results with a dataset containing older orders

```
'-> Limit: 10 row(s)

-> Sort: TotalUnitsSold DESC, p.Title, limit input to 10 row(s) per chunk

-> Table scan on <temporary>

-> Aggregate using temporary table

-> Nested loop inner join (cost=7.46 rows=8)

-> Nested loop inner join (cost=4.66 rows=8)

-> Filter: (o.OrderDate >= <cache>((curdate() - interval 30 day)))
(cost=1.86 rows=8)

-> Covering index range scan on o using
idx_orders_orderdate_orderid over ('2025-11-04 00:00:00' <= OrderDate)
(cost=1.86 rows=8)

-> Index lookup on op using PRIMARY (OrderID = o.OrderID)
(cost=0.263 rows=1)

-> Single-row index lookup on p using PRIMARY (ProductID =
op.ProductID) (cost=0.263 rows=1)

'

• Some orders fell outside the 30-day window.

• EXPLAIN showed a covering index range scan on the new index.

• Rows examined dropped from the full table (14) to only the recent orders ( $\approx 8$ ).

• Joins on Order_Product and Product continued to use efficient primary-key
lookups.
```